

Figure 1: Compared to CycleAdapt method, our predictions show more aligned semantics (*i.e.*, climbing-the-stairs and walking) with the ground truth (GT).

References

Stathopoulos, A., Han, L., and Metaxas, D. Score-guided diffusion for 3d human recovery. In *CVPR*, pp. 906–915, 2024.

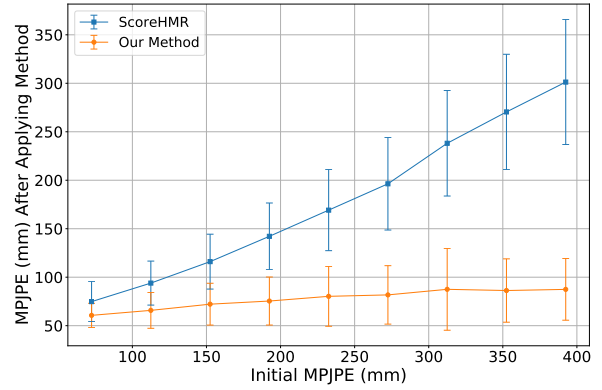


Figure 2: Methods that optimize 3D pose outputs without model fine-tuning are heavily dependent on the quality of initial predictions. We demonstrate this limitation by comparing with ScoreHMR (Stathopoulos et al., 2024), which shows significant sensitivity to poor initial predictions. In contrast, ours maintains robust performance.